

PokerCRR

Applying Cox-Ross-Rubinstein binomial pricing logic to poker decision-making

This project treats a poker hand as a discrete-time stochastic process, the same structure a CRR binomial tree uses to price an option: equity plays the role of the underlying's price, streets play the role of time steps, and betting decisions are priced using the same no-arbitrage logic used for options and other derivatives. Four modules implement this end to end: Monte Carlo equity estimation, martingale-based bet pricing, a Kuhn Poker GTO solver trained via Counterfactual Regret Minimization, and Kelly Criterion bankroll sizing.

Contents

Methodology: the CRR-to-poker mapping	2
1. Equity Estimation and Pricing	2
2. Bankroll Sizing (Kelly Criterion)	3
3. GTO Solving via Counterfactual Regret Minimization	4
4. Extended Validation	5
4.1 Exact validation against Kuhn Poker's closed-form solution	5
4.2 Backtest across many hands	6
4.3 Testing the Kelly module's independence assumption	6
4.4 Extending the GTO solver: Leduc Hold'em	7
Limitations and Future Work	7
Appendix: Core Pricing Logic	8
Summary	8


Methodology: the CRR-to-poker mapping

The Cox-Ross-Rubinstein (CRR) model prices an option by building a binomial tree: at each discrete time step, the underlying's price moves up or down, and the option's value is computed by working backward from the known payoffs at expiry, applying risk-neutral probabilities at every node. A poker hand has the same discrete-time structure: it moves through a fixed sequence of steps (preflop, flop, turn, river), and at each step new public information (a card) shifts the probability of winning, exactly as a price move shifts an option's moneyness.

Options / CRR concept	Poker analogue	Implemented in
Underlying's price	Equity (win probability)	equity/monte_carlo.py
Risk-neutral pricing	Pot odds / EV of a call	pricing/martingale.py
Backward induction	Solving for equilibrium play	gto/kuhn_solver.py
Position sizing under a model	Bankroll sizing given an edge	bankroll/kelly.py

1. Equity Estimation and Pricing

The equity module runs Monte Carlo simulation (20,000 random runouts per hand) to estimate win probability at any point in a hand. The pricing module then applies pot-odds logic, the poker equivalent of risk-neutral pricing, to decide whether a call is +EV. Below is a broader set of six hands spanning strong made hands, draws, and air, all facing the same bet size (pot=100, facing a 50-chip bet).




Two representative cases from the table below, shown visually: a strong overpair clears the pot-odds threshold comfortably, while pure air does not.

Hand	Equity	EV of call	Decision	Full Kelly
Overpair (AA)	85.8%	+121.7	CALL	78.7%
Overpair (KK)	83.7%	+117.4	CALL	75.6%
Nut flush draw	72.8%	+95.6	CALL	59.2%
Top pair, weak kicker	85.4%	+120.9	CALL	78.2%
Gutshot draw	40.6%	+31.1	CALL	10.8%
Air (72o)	22.1%	-5.7	FOLD	0.0%

Note the gutshot draw: at 40.6% raw equity it looks strong at first glance, but this is a single-street snapshot -- the module prices the immediate call, not the full multi-street decision tree, which is a deliberate scope limitation discussed below.

2. Bankroll Sizing (Kelly Criterion)

Once a decision is +EV, the Kelly Criterion converts the edge into an optimal stake size as a fraction of bankroll, maximizing long-run geometric growth without risking ruin. Full Kelly is theoretically growth-optimal but highly volatile in small samples; the module also supports fractional Kelly (e.g. half-Kelly) for lower variance at a modest cost to long-run growth, a trade-off widely used in both poker bankroll management and position sizing in trading.




Optimal Kelly stake as equity varies, for a fixed pot-size bet. Below 50% equity (the break-even threshold for this bet size) the optimal stake is exactly zero -- Kelly correctly refuses -EV bets rather than sizing them down. A real bug was found and fixed during development here: an earlier version double-counted the returned stake in the payoff calculation, shifting the break-even point to 33% instead of the correct 50% -- caught precisely because the chart didn't match the known theory.

3. GTO Solving via Counterfactual Regret Minimization


Kuhn Poker (a simplified 3-card game: Jack, Queen, King, one betting round) has a known, provably optimal equilibrium strategy, famous for the result that a rational player should bluff the worst hand (the Jack) with probability in $[0, 1/3]$. Rather than hard-coding this answer, the solver learns it from 200,000 iterations of self-play using Counterfactual Regret Minimization (CFR), the same self-play logic used by modern superhuman poker AI (e.g. Libratus, Pluribus), applied here at a small enough scale to solve exactly and verify against the textbook result.

How CFR works, step by step:

- 1 Play the game against itself, tracking every decision point (an 'information set': a player's own card plus the betting history so far).
- 2 At each information set, maintain a running 'regret' for every action -- how much better that action would have performed versus the strategy actually played.
- 3 Choose actions in proportion to their positive regret ('regret matching'): actions that would have done better get played more often in future iterations.
- 4 Track a running average of the strategy played at each information set over all iterations -- this average, not the strategy at any single iteration, is what converges to the Nash equilibrium.
- 5 Repeat for many iterations; the average strategy's exploitability provably shrinks toward zero as iterations increase.



The solver's estimated game value converges to the known theoretical value of $-1/18$ as training iterations increase, confirming the self-play process is actually finding the equilibrium rather than an arbitrary fixed point.



Learned bet/call frequency at every information set. The model independently recovers the $\sim 1/3$ opening bluff frequency with the Jack, and the theoretically required $1/3$ bet frequency with a Queen after the opponent checks then bets -- neither value was hard-coded; both emerged purely from repeated self-play.


4. Extended Validation

The original four modules were each checked against a known benchmark in isolation. This section goes further: an exact closed-form comparison (not just a converged game value), a multi-hand backtest instead of isolated examples, an explicit test of the Kelly module's independence assumption, and an extension of the GTO solver to a larger game with two betting rounds.

4.1 Exact validation against Kuhn Poker's closed-form solution


Kuhn (1950) proved a closed-form family of equilibria for this game, parametrized by a single value (the Jack's bluff frequency). Rather than only checking that the CFR solver's game value is close to the known $-1/18$, every one of the 12 information sets was compared directly against its exact theoretical probability. Mean absolute error across all 12 was 0.16%, with a maximum error of 1.0% at any single decision point -- a considerably stronger test than matching a single aggregate number.

4.2 Backtest across many hands



Simulated bankroll (log scale) over 300 independent random hands, following the pipeline's call/fold decisions and staking quarter-Kelly on each +EV call. The steady upward slope on a log scale is the signature of geometric growth: the per-hand growth rate was +2.4%, compounding into a large total return over 300 hands -- a reminder that Kelly-style compounding, not any single decision, drives the aggregate outcome.


4.3 Testing the Kelly module's independence assumption



The base Kelly module sizes every decision as if it were independent, a limitation flagged earlier in this report. To test the cost of that assumption directly, 2,000 sessions of 15 hands were simulated where hands within a session share a common regime (correlated, not independent) -- and two staking approaches were compared: naive Kelly (ignores the correlation) and a half-sized 'shrunk' Kelly (a standard practitioner response to correlated risk). Naive Kelly grew the average bankroll faster (+73.8% vs. +27.5%) but produced meaningfully worse tail risk -- an average max drawdown of 26.9% versus 14.8%, and a worst observed drawdown of 94.4% versus 72.6%. The chart shows the full distribution: shrunk Kelly's drawdowns cluster near zero, while naive Kelly has a long tail of severe drawdowns.

4.4 Extending the GTO solver: Leduc Hold'em

Kuhn Poker's GTO solver was extended to Leduc Hold'em: 6 cards, two betting rounds, and a public card dealt between them -- the smallest standard benchmark game with an actual 'flop'. Pot contributions are tracked exactly through the recursion (not inferred from the action history), so payoffs are pot-accurate. Leduc has no simple closed-form solution the way Kuhn does, so it is validated differently: by checking that the estimated game value stabilizes as training increases, the standard convergence signature for a solver with no known closed form.



Estimated game value for Player 0 across increasing training iterations. The value stabilizes around -0.12 after roughly 50,000 iterations, indicating the solver has converged to a stable strategy rather than continuing to drift.

Limitations and Future Work

Section 4 directly addressed two of the limitations from the original four modules: the GTO solver was extended beyond Kuhn Poker (Leduc Hold'em), and the Kelly module's independence assumption was tested explicitly. What remains open:

- Even Leduc Hold'em is far smaller than real Texas Hold'em, which has a vastly larger information-set space. Scaling further would require abstraction techniques (card bucketing, action abstraction) used in production-grade solvers, not brute-force CFR -- Leduc is a meaningful step up from Kuhn, not a solution to the full-scale problem.
- The pricing module, and the backtest built on it, still evaluate one decision point at a time rather than pricing the entire remaining game tree (all future streets and bet sizes) -- the poker equivalent of pricing a path-dependent option rather than a single-period one.
- Equity estimation assumes uniformly random opponent hands. A more realistic model would condition on a plausible opponent range inferred from the betting action so far.
- The multi-period Kelly comparison in Section 4.3 uses a simplified, single-factor model of within-session correlation (one shared regime per session). A fuller treatment would estimate correlation structure empirically rather than assuming it, closer to how a real portfolio's covariance matrix is estimated rather than assumed.

Appendix: Core Pricing Logic

The two functions below are the heart of the pricing module -- pot odds (the break-even equity threshold) and the fair, zero-EV bet size at a given equity, the poker analogue of a no-arbitrage price.

```
def pot_odds(bet, pot):
    """Required equity to break even on a call."""
    return bet / (pot + bet)

def fair_bet_size(pot, equity):
    """
    Bet size at which calling has exactly zero EV, given equity.
    Solved from: equity*(pot+bet) = (1-equity)*bet
    """
    if equity >= 0.5:
        return float('inf')
    return equity * pot / (1 - 2 * equity)
```

The CFR regret-matching update, the mechanism that lets the GTO solver converge to equilibrium purely through self-play rather than a hand-coded solution:

```
def get_strategy(self, realization_weight):
    strategy = [max(r, 0.0) for r in self.regret_sum]
    total = sum(strategy)
    if total > 0:
        strategy = [s / total for s in strategy]
    else:
        strategy = [1.0 / NUM_ACTIONS] * NUM_ACTIONS
    for a in range(NUM_ACTIONS):
        self.strategy_sum[a] += realization_weight * strategy[a]
    return strategy
```

Summary

All four core modules were built, tested against known theoretical benchmarks (AA preflop equity ~85%, Kuhn Poker game value -1/18, classic Kelly break-even points), and connected into a single decision pipeline. The extended validation in Section 4 pushed this further: exact node-by-node agreement with Kuhn Poker's closed-form solution, a multi-hand backtest instead of isolated examples, a direct test of where the Kelly module's independence assumption breaks down under correlated decisions, and an extension of the GTO solver to a two-round game with a public card. One real bug was found and fixed during development using the visualizations themselves as a validation tool, not just a presentation layer. The project's contribution is not any individual algorithm, each of which is well documented in the literature, but the explicit framing of poker decisions through the language of no-arbitrage derivatives pricing: equity as an underlying's price, bets as options to be priced, and bankroll sizing as position sizing under a known edge.